

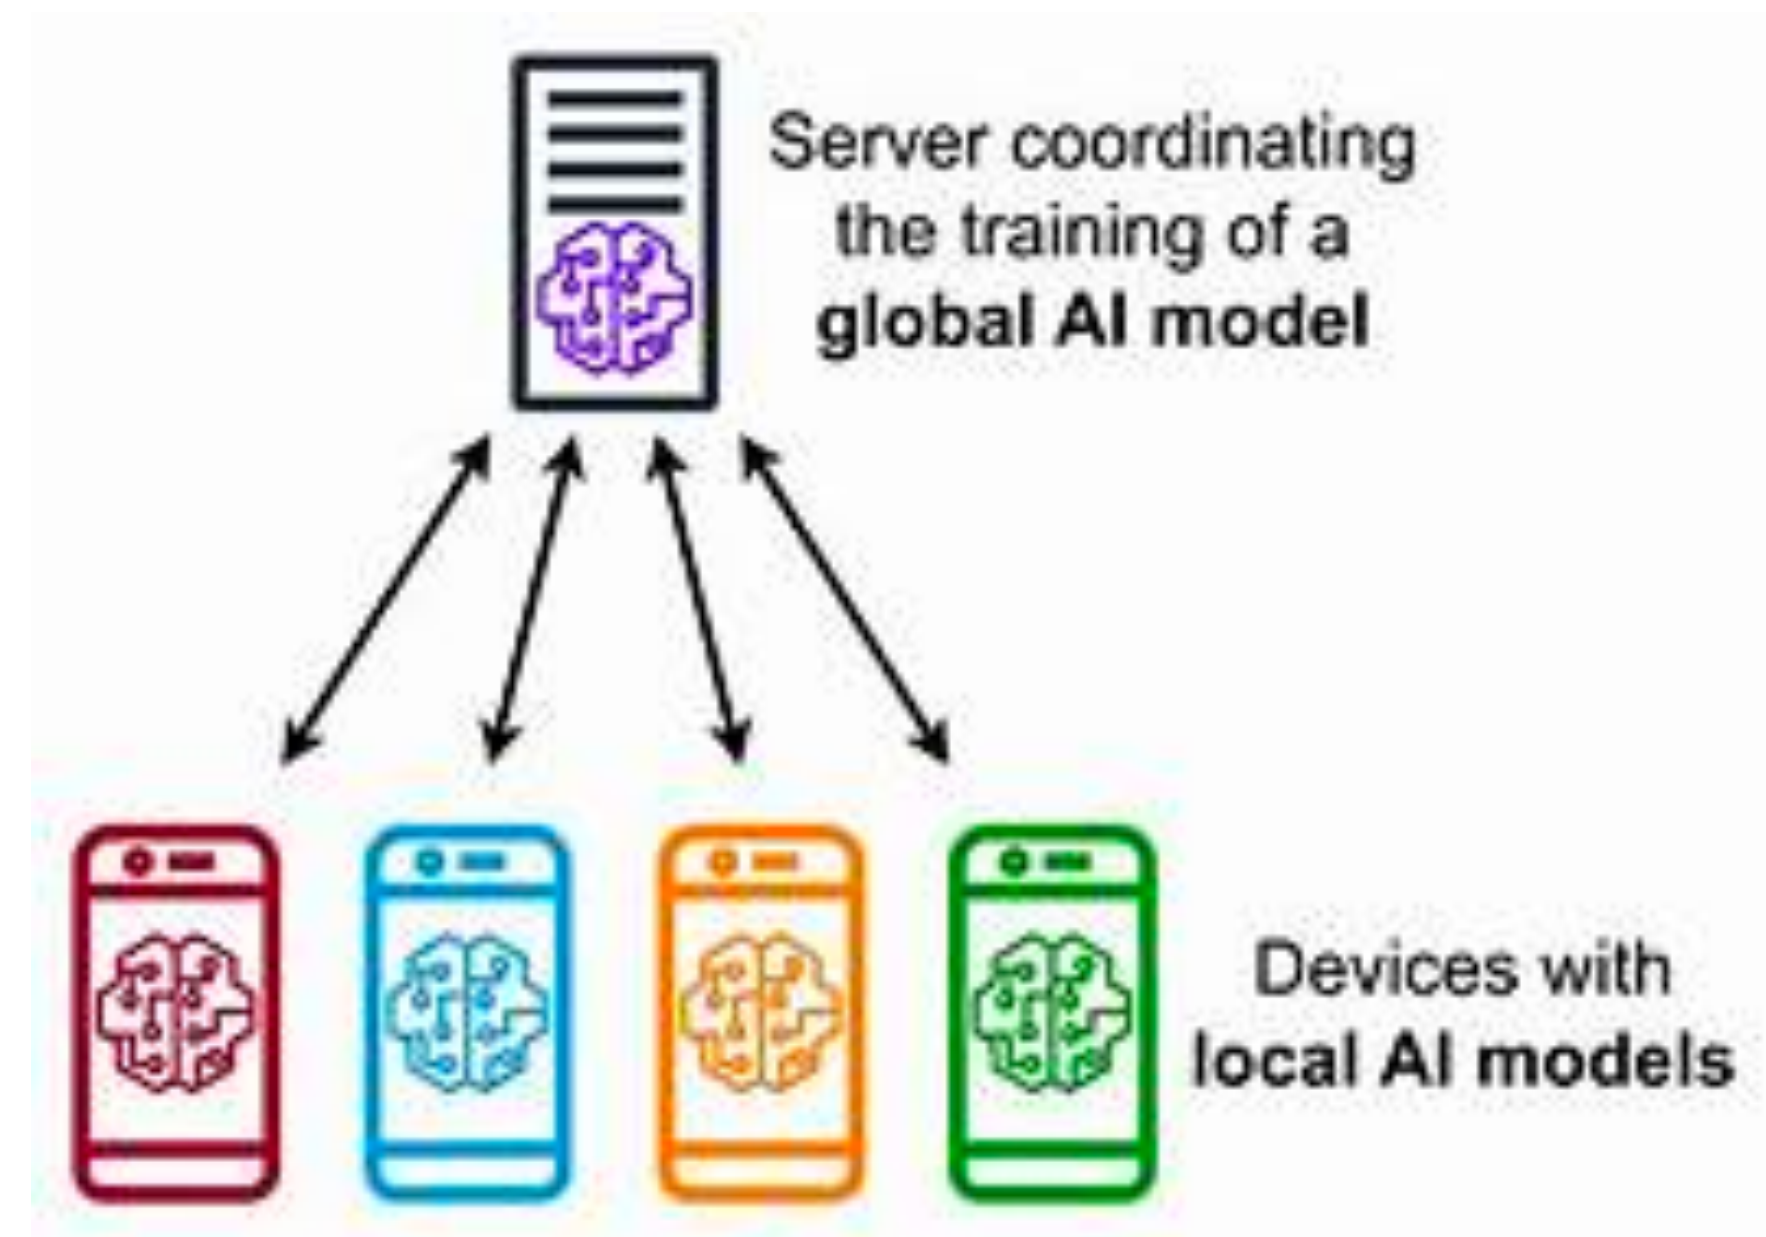
Federated Learning for Privacy- Preserving Machine Learning

Final Project Presentation

Raj Kumar Maurya, 2024.05.02

Introduction

- Federated Learning (FL) is a decentralized machine learning approach where model training is performed across multiple edge devices or servers holding local data, without exchanging them. This project outlines the implementation of Federated Learning in a real-world scenario.



Federated Learning Algorithms

- **Federated Averaging (FedAvg):** It is one of the foundational federated learning algorithms. In FedAvg, each participating device trains its local model on its own data, and then the model updates are aggregated (averaged) at a central server. This process helps in preserving data privacy since raw data doesn't leave the devices.
- **Federated Stochastic Gradient Descent (FedSGD):** Similar to traditional SGD, FedSGD optimizes a global model by updating parameters using gradients computed locally on each device. These local gradients are then aggregated to update the global model.
- There are many other algorithms:
 - Federated Momentum (FedMo)
 - Federated Proximal (FedProx)
 - Federated Learning with Adaptive Gradient Clipping (FLAG)
 - Federated Learning with Differential Privacy (FLDP)
 - Federated Learning with Differential Privacy (FedDP)
 - Federated Learning with Secure Aggregation (FedSecAgg)
 - Federated Learning with Homomorphic Encryption (FedHomEnc)

Objective

- **Understand** Federated Learning: Study the principles, challenges, and benefits of federated learning.
- **Implement** Federated Learning techniques to leverage distributed data for model training.
- **Demonstrate** the feasibility of Federated Learning in improving model performance while maintaining data privacy.
- **Apply** federated learning algorithms to a specific use case and do comparative analysis

Methodology

- **Literature Review:** Survey existing research on federated learning
- **System Architecture:** Design the system to implement the federated learning in a simulated environment or actually distributed systems
- **Algorithm Selection:** Choose appropriate federated learning algorithms based on use case requirements.
- **Implementation:** Develop a proof-of-concept system. Dataset and the problem statement to be finalised.
- **Evaluation Metrics:** Define metrics to assess model performance, communication efficiency, and privacy preservation.

FedAvg vs FedSGD

Aspect	FedSGD	FedAvg
Aggregation Method	Aggregates gradients	Aggregates model parameters
Update Strategy	Updates global model directly with gradients	Averages model parameters to update global model
Communication Overhead	Higher due to gradient transmission	Lower due to parameter transmission
Privacy Preservation	Gradients may contain more information about local data	Averaged parameters may preserve privacy better
Frequency of Synchronization	Every iteration	After several local updates

Expected Outcomes

- **Functional Prototype:** A working federated learning system demonstrating privacy-preserving model training.
- **Performance Evaluation:** Comparative analysis of the two federated learning algorithms chosen.
- **Documentation:** Detailed project report, code snippets, and experimental results.

Feedbacks from 1st Presentation

- *Focus on the distributed aspects*
- *Try with smaller models*
- *Use cluster of VMs provided instead of simulating locally*
 - *Since VM was not accessible over VPN, went ahead with simulation on multi-core system after getting confirmation from the TA*

Solution Overview & Summary

- I have implement FedAvg and FedSGD Federated Learning technique on a **synthetic data** for a **regression problem** using **2 Layers Vanilla Neural Network** that highlights the properties of these algorithms.
- I have ran **various experiments** and solved multiple scenarios of the problems.
- I have attached all the files in the shared repo, it contains **REPORT.md** which contains the project report.
- I have added **README.md** that contains steps to setup and run the project.
- **Topics** covered ahead:
 - Data Preparation
 - Training
 - Metrics
 - Distributed System Perspective Analysis
 - Raw Outputs

Data Preparation

- Used dataset from a file on each node based on row number where **row % rank_of_the_process == 0**
- **Synthetic Data** is **polynomial** in nature with some white gaussian noise with a constant mean and standard deviation. Data has been prepared in a way so that it can illustrate the benefits of federated learning even on very short samples and tiny model.
- I have used entire data as **test data** on each process to maintain consistency in evaluation.
- I have plotted the **data distribution** and put it in images folder for your reference.
- **Privacy Preserving Feature:** Since each node has it's own training data, data is not shared between the systems which preserves privacy. Here I have simulated this using the a file reading based on row ids, but it illustrates the privacy idea.

Training

- In order to **benchmark** the performance, I implemented the same architecture (that I used in federated learning) neural network (NN) with any distributed federated learning to get a **baseline**.
- Developed multiple **variants** of FedAvg method by changing the point of aggregation in various experiments. The **global aggregation** turned out to be best on this specific data and model architecture, which is might not be true in other cases.
- Also, implemented the FedSGD technique where we were doing the **aggregations** at every epoch of the training. Also, tried variation with **synchronisation** after every 10 epochs, which was sub-optimal.
- Trained the models in a way that the results are **reproducible**.

Metrics

- For the given data and model, **FedSGD** was **closer** to the baseline than FedAvg technique. It's not conclusive to say FedSGD is better than FedAvg as it **depends on the usecase, model, data and scenario**.
- For **Baseline** Model:
 - Final loss: 1141.0064697265625
 - R-squared: 0.7502739980476314
 - Mean Absolute Error: 27.496908950417136
- For **FedAvg**:
 - After Reducing at Rank 0, Final loss: 1318.7928466796875
 - After Reducing at Rank 0, R-squared: 0.7113628476964176
 - After Reducing at Rank 0, Mean Absolute Error: 28.933150488261308
- For **FedSGD**:
 - After Reducing at Rank 2, Final loss: 1104.7113037109375
 - After Reducing at Rank 2, R-squared: 0.7582177084580953
 - After Reducing at Rank 2, Mean Absolute Error: 27.069782214178403

Distributed System Perspective Analysis

- **Message Size:**
 - From total number of messages/parameters perspective **per synchronisation**, the number of messages remain same, but since **FedSGD synchronises more frequently** than FedAvg, overall it sends more number of messages in total.
 - Moreover, in practice, since storing **gradients require higher floating precision points** than parameters, FedSGD has even more requirements in terms of overall **message size** capacity than FedAvg.
- **Communication:**
 - **FedSGD** requires **more bandwidth** than FedAvg as it synchronises after every gradient computation while FedAvg runs after a few epochs.
- **Deadlock:**
 - I have used **Reduceall** API of MPI which is a **blocking** API and since there is **no interdependency** anywhere else, deadlock will not occur in this scenario.

Raw Outputs

Refer to the DEMO and the REPORT.md

Aspect	FedSGD	FedAvg
Aggregation Method	Aggregates gradients	Aggregates model parameters
Update Strategy	Updates global model directly with gradients	Averages model parameters to update global model
Communication Overhead	Higher due to gradient transmission	Lower due to parameter transmission
Privacy Preservation	Gradients may contain more information about local data	Averaged parameters may preserve privacy better
Frequency of Synchronization	Every iteration	After several local updates

2nd Presentation Comments

Add these to project code and share the updated code today

- Plots
 - Error vs epochs for all the algos
 - Message size plot for fed algo [cumulative messages vs epochs]
 - Number of message [cumulative #msg vs epochs]

2nd Presentation Feedback for Future Works

- Implement the algorithms for more realistic scenarios:
 - Async
 - Non-iid
 - Implement on actual distributed systems instead of simulation
- Distributed GPU training
 - NCCL (Nvidia's Collective Commons Library on top of MPI)
 - CUDA
- Connect with Manaswi ma'am for the above

Thank you!

Q&A